Method	CU (%) ↑	UA (%) ↑	ASR (%) ↓	Total LLM Calls ↓
No Defense	100.0	100.0	61.3	406
PIGuard	100.0	100.0	61.3	410
ProtectAI	68.1	68.5	58.9	416
Progent	100.0	100.0	61.3	1025
DRIFT	97.2	99.2	60.5	2384
DAFR	93.1	99.2	**0.0**	690

Table 1: Method comparison on the six-domain synthetic tool-agent dataset with DeepSeek Flash. All configurations use the same 196 tasks, tools, and execution budgets. CU and UA measure original-task completion on 72 clean and 124 attacked tasks, respectively. ASR measures attainment of the injected redirection or verification-bypass objective on attacked tasks. Utility and attack success are evaluated independently. Rates use pooled task counts. Total LLM Calls counts planner and defense LLM calls across all 196 tasks.