Method	Mean ± SD (s)	Median (s)	LLM Calls	Episodes
No Defense	12.30 ± 1.62	11.71	6.00	4
PIGuard	15.11 ± 3.25	15.44	7.00	4
ProtectAI	16.36 ± 2.19	16.72	7.25	4
Progent	47.09 ± 18.07	48.00	19.00	4
DRIFT	34.89 ± 10.08	34.07	19.00	4
DAFR	15.88 ± 2.77	16.59	7.50	4

Matched sequential pilot: two paired tasks, four executions per method. Includes agent, defense and tools.
Model loading and local detector warm-up excluded; measured after the chain experiment finished.